

Reading the Linearizability paper of Herlihy and Wing

Uri Abraham *

May 28, 2013

Abstract

We read Herlihy and Wing important paper on the concept of linearizability, and reconstruct that text in the Tarskian event structure paradigm. With linearizability defined in terms of Tarskian system executions we obtain a different correctness proof for the queue algorithm that Herlihy and Wing present in that paper.

1 Definition of linearizability by means of Tarskian system executions

The definition that we give here for *linearizability* is virtually the same as the definition that was originally given by Herlihy and Wing in their very influential paper¹, but since we use Tarskian system executions rather than histories in our definition, the two definitions may superficially appear to be different. We will argue in this article that (with some evident exaggeration) the authors of the “Linearizability” paper have spoken in Tarskian prose even while presenting their linearizability theory in terms of states and histories.

Recall the definition of moment based system execution languages. Such a language contains the sort *Moment*, and there are two functions *Left_End*, and *Right_End* so that, for any event e , $[Left_End(e), Right_End(e)]$ is the temporal extension of e (an interval in the linearly ordered set *Moment*). Assume a class \mathcal{N} of system executions that interpret L so that in every

*Departments of Mathematics and Computer Science, Ben-Gurion University. Mathematical Models for Concurrency course, 2010.

¹Maurice Herlihy, Jeannette M. Wing: Linearizability: A Correctness Condition for Concurrent Objects. ACM Trans. Program. Lang. Syst. 12(3): 463-492 (1990)

$N \in \mathcal{N}$ the events are linearly ordered by the precedence relation $<^N$ of that structure. An arbitrary system execution M that interprets L (for which $<^M$ is not necessarily a linear ordering) is *linearizable* (with respect to \mathcal{N}) if there is a linear ordering $<^*$ of the set of events of M so that $<^*$ extends $<^M$ and so that if we replace $<^M$ with $<^*$ in M then the resulting structure N is in \mathcal{N} .

The following example should clarify this notion of linearizability. Recall the language L_{queue} which we used to specify serial queues. There were two sorts *Event* and *Data*, two predicates *enq* and *deq*, a function *Val*, and most importantly the function γ that yielded for every successful dequeue event d its corresponding enqueue event $\gamma(d)$. (There were additional items in L_{queue} but these were the most important). Recall also that we defined a sentence in that language (let's call it ψ) which specifies the behavior of serial queues. The main ingredients of this statement were that the enqueue and dequeue events are linearly ordered by $<$, that $Val(d) = Val(\gamma(d))$ for any dequeue event d (unless d returns the “empty” value), that $\gamma(d) < d$, and that γ is order preserving and one-to-one (which is a reflection of the FIFO property). Let \mathcal{N} be the set of all system executions N that interpret L_{queue} and which satisfy ψ . That is, \mathcal{N} is the set of all system executions in which the queue operations correctly represent a serial FIFO regime. Then an arbitrary system execution interpretation of L_{queue} is linearizable if its partial ordering of events can be augmented to yield a structure in \mathcal{N} .

2 The Herlihy and Wing queue algorithm

The Herlihy and Wing queue algorithm is in figure 1. The algorithm implements a distributed queue that allows Enqueue and Dequeue operations by concurrently operating processes.

The algorithm assumes an infinite array *items* indexed with natural numbers: *items*[1], *items*[2], ..., and initially all entries *items*[i] have the “empty” value $*$ (which is not a *Data* value). Every cell *items*[i] is a serial register. We also assume a global variable *back* of type natural number which is initially 1. The role of *back* is to hold the index of the first cell in *items* that is available for the next *Enqueue* operation. Operation *Enqueue*(x) (where x is the *Data* value with which the operation is invoked) begins with an atomic action $i := INC(back)$ which assigns to i the value of *back* and then increases *back* by one. The algorithm would not work properly if this action is split into separate assign and increase parts, but we do assume an atomic operation which ensures that operation *INC* is executed instantly.

<p><i>Enqueue</i>(x)</p> <ol style="list-style-type: none"> 1. $i := INC(back)$; 2. $items[i] := x$. 	<p><i>Dequeue</i></p> <p>while true</p> <ol style="list-style-type: none"> 1. $range := back - 1$; 2. for $i := 1$ to $range$ do $x := swap(items[i], *)$; if $x \neq *$ then return x
---	--

Figure 1: The Hrlhy and Wing queue protocol.

The Dequeue operation consists of one or more “cycles”. A cycle is an execution of lines 1 and 2. A successful cycle is a cycle that returns. So a Dequeue operation may consist of a finite number of cycles, the last being successful, or of an infinite number of non-successful cycles. In executing line 1 of a cycle, variable $back$ is read and the value $range - 1$ is assigned to local variable $range$ of the executing process. The $swap$ operation in line 2 is an atomic operation that reads the value of $items[i]$, assigns the value obtained into local variable x , and replaces this value with $*$. This is an atomic action.

We shall examine now an arbitrary execution of this algorithm, and state some properties that are rather easy to derive informally. These properties are collected in Figure 2, and we will prove that they imply that the queue is linearizable. The properties refer to high-level events, and so we first define these events. An Enqueue operation is an execution of lines 1 and 2 of the *Enqueue* code by one of the concurrent processes. If E is an Enqueue operation then $inc(E)$ is the execution of line 1 in E . $Val(inc(E))$ is the value of i that is obtained in this action. Then $asn(E)$ is the execution of line 2, in which $x = Val(E)$ is assigned to $items[Val(inc(E))]$.

There are two possibilities for a Dequeue operation execution D . D is not successful if it is non-terminating, and in this case it contains an infinite number of cycles. When D is terminating only its last cycle C is important for the proof. We denote with $b(D)$ the read action in C of the $back$ variable, and we let $r = Val(b(D))$ be the value obtained in this read. Then $s(D)$ denotes the last swap action in C , which is the action in which the “empty” token $*$ is written on $items[r]$ instead of the *Data* value $x = Val(D)$ which is returned.

Since a Dequeue operation, even when successful, may contain unsuccessful cycles that are not relevant for the proof we define *Deq* events which

are just the successful cycles of Dequeue events. So if D is any Dequeue event then its last cycle is a *Deq* event. When D is a *Deq* event then $b(D)$ and $s(D)$ denote the read of *back* and the swap actions in D , so that $D = \{b(D), s(D)\}$. And $begin(D)$, $end(D)$ are the Moments at which $b(D)$ and $s(D)$ are executed.

We also define *Enq* events, which are just the Enqueue operation executions. So any *Enq* event E has the form $E = \{inc(E), asn(E)\}$. (And $begin(E)$, $end(E)$ are their Moment values.)

We define the function γ on the terminating *Deq* events. If D is a terminating *Deq* event and $s(D)$ its (last and successful) swap action executed on $items[i]$, then the *Data* value x obtained must have been assigned to $items[i]$ by some *Enq* operation E in executing the write operation $asn(E)$. Note that E is uniquely determined by index i since different *Enq* events access different indexes. We then define $E = \gamma(D)$.

The *HW* properties enumerated in Figure 2 are expressed in the L_{queue} language. We use the letter D to denote *Deq* events (which are successful), and E refers to *Enq* events. Since the domain of γ is the set of *Deq* events, when we write $\gamma(X)$ it is implicitly evident that $Deq(X)$ holds.

In class we gave intuitive arguments to explain why the *HW* properties hold in every execution of the algorithm. Our main aim is to show that they imply linearizability. This is done in the following section ???. Only after seeing the relevance of these properties we shall return and prove more formally that they hold.

2.1 High level properties and their linearizability consequence

Our aim is to prove that in any Tarskian system execution in which the *HW* properties 1–3 hold, a linearizable queue is implemented.

Theorem 2.1 *In the context of Moment Based System Executions, properties HW imply that there exists a linear ordering $<^*$ that extends $<$ and which satisfies the following.*

1. $\gamma(D) <^* D$ for every *Deq* event D .
2. The function γ is $<^*$ -order preserving. That is, for every successful *Deq* events D_1, D_2 : $D_1 <^* D_2$ iff $\gamma(D_1) <^* \gamma(D_2)$.

A main idea for the proof is to linearly order the *Deq* events by the ordering of their end-points. That is, we intend to define $D_1 <^* D_2$ iff $end(D_1) < end(D_2)$. This decision immediately entails the ordering of the

1. Every event is either a *Deq* or else an *Enq* event. All *Enq* events are terminating. If D is a terminating *Deq* event then $\gamma(D)$ is defined, it is an *Enq* event and $Val(\gamma(D)) = Val(D)$.
2. The functions *begin*, *end* are defined on the events and yield Moment values. These are one-to-one functions. The γ function is also one-to-one.
3. If $E = \gamma(D)$ then $Val(E) = Val(D)$ and

$$begin(E) < begin(D), \text{ and } end(E) < end(D). \quad (1)$$

4. If $E_0 < \gamma(D)$ (where E_0 is an *Enq* event) then, for some *Deq* event D_0 , $E_0 = \gamma(D_0)$ and $end(D_0) < end(D)$.

A stronger form holds: If $E = \gamma(D)$ and E_0 is an *Enq* event so that $E_0 < D$ and $begin(E_0) < begin(E)$, then for some *Deq* event D_0

$$E_0 = \gamma(D_0) \text{ and } end(D_0) < end(D).$$

Figure 2: The HW list of properties of the high level events of the Herlihy and Wing queue algorithm.

corresponding *Enq* events. That is, we shall request that for *Enq* events E_1 and E_2 of the form $E_\ell = \gamma(D_\ell)$ (for $\ell = 1, 2$) the equivalence $E_1 <^* E_2$ iff $\text{end}(D_1) < \text{end}(D_2)$ holds. In addition, we shall demand that $\gamma(D) <^* D$ for every successful *Deq* event D . And, of course, we demand that the relation $<^*$ extends the given partial ordering $<$. Combining these demands we get a relation \prec on the *Enq/Deq* events which is not necessarily an ordering relation. The main mission in the proof of Theorem 2.1 is the proof that \prec contains no cycles and hence that the transitive closure of \prec is indeed a partial ordering. This is achieved in a series of lemmas.

We are going to define five relations on the events, \prec_i for $i = 1, \dots, 5$, and then define \prec as their union. That is $A \prec B$ iff $A \prec_i B$ for some $i = 1, \dots, 5$. The main step in the proof is showing that \prec has no cycles and hence that its transitive closure \prec^* is an irreflexive relation (and thus a partial ordering). Now \prec^* is not necessarily a linear ordering, but a simple extension of \prec^* yields the desired linearization.

Define the following relations:

1. For any events A and B : $A \prec_1 B$ iff $A < B$.
2. For *Enq* and *Deq* events E and D define $E \prec_2 D$ iff $E = \gamma(D)$.
3. For *Deq* events D_1 and D_2 define $D_1 \prec_3 D_2$ iff $\text{end}(D_1) < \text{end}(D_2)$.
4. For *Enq* events E_1 and E_2 define $E_1 \prec_4 E_2$ iff for some *Deq* events D_1 and D_2 we have that $E_1 = \gamma(D_1)$, $E_2 = \gamma(D_2)$, and $\text{end}(D_1) < \text{end}(D_2)$.
5. For *Enq* events E_1 and E_2 define $E_1 \prec_5 E_2$ iff E_1 is in the range of γ , but E_2 is not.

Note that these relations are not disjoint and $<$ may have a non-empty intersection with any of the other relations. For example, if D_1, D_2 are *Deq* events and $D_1 < D_2$ then also $D_1 \prec_3 D_2$. Relations \prec_i for $i > 1$ however, are mutually exclusive.

Define $X \prec Y$ iff $X \prec_i Y$ for some $i = 1, \dots, 5$. Let \prec^* be the transitive closure of \prec . That is, $X \prec^* Y$ iff there exists a sequence X_0, \dots, X_n of events (with $n > 0$) where $X_0 = X$ and $X_n = Y$ and so that $X_k \prec X_{k+1}$ for every $k < n$. The following sequence of lemmas will finally prove that \prec^* is an irreflexive relation.

Lemma 2.2 *If D_1, D_2 are Deq events then $D_1 \prec D_2$ iff $D_1 \prec_3 D_2$*

Proof. $D_1 \prec_1 D_2$ implies $D_1 \prec_3 D_2$, and relations \prec_2 and \prec_4 are not applicable. **q.e.d**

Note that \prec_2 is a transitive relation simply because it is never the case that $X \prec_2 Y \prec_2 Z$ (which would entail that Y is both an *ENQ* and a *DEQ* event. Relation \prec_3 is also transitive (because γ is one-to-one). Similarly, all other \prec_i relations are transitive, and the following is easy to prove.

Lemma 2.3 *Each \prec_i for $i = 1, \dots, 5$ is an irreflexive and transitive relation. Relation \prec_3 is in fact a linear ordering of the terminating *Deq* events which extends the $<$ relation over these events.*

Lemma 2.4 *If $E_0 \prec E_1$ are *Enq* events and E_1 is in the range of γ then $E_0 \prec_4 E_1$ (and in particular E_0 is also in the range of γ).*

Proof. Suppose E_0 and E_1 are both *Enq* events, $E_0 \prec E_1$ and $E_1 = \gamma(D_1)$. Relations \prec_2 and \prec_3 clearly do not apply to $E_0 \prec E_1$ since these relations never relate two *Enq* events. But $E_0 \prec_5 E_1$ is also impossible since E_1 is in the range of γ . This leaves two possibilities: $E_0 < E_1$ and $E_0 \prec_4 E_1$. To prove the lemma we show that $E_0 < E_1$ implies $E_0 \prec_4 E_1$. Since $E_1 = \gamma(D_1)$, property 3 of the HW list implies (in case $E_0 < E_1$) that there exists some *Deq* event D_0 with $E_0 = \gamma(D_0)$ and $\text{end}(D_0) < \text{end}(D_1)$. That is $E_0 \prec_4 E_1$ as required. **q.e.d**

Lemma 2.5 *The \prec relation has no cycles of length two.*

Proof. Since \prec is irreflexive, it is never the case that $X \prec X$. So we have to show that there are no two events $X \neq Y$ such that $X \prec_i Y \prec_j X$ for some i and j . To prove this let's check the different possibilities. First, $i = j$ is impossible since \prec_i is transitive, and $X \prec_i Y \prec_i X$ would contradict the irreflexivity of \prec_i . So it must be that $X \prec_i Y \prec_j X$ where $i \neq j$.

1. Suppose both X and Y are both *Deq* events. We noted already (Lemma 2.2) that if D_1, D_2 are *Deq* events then $D_1 \prec D_2$ implies that $D_1 \prec_3 D_2$. Hence $X \prec Y \prec X$ would bring us back to $X \prec_3 Y \prec_3 X$ which we ruled out (by transitivity and irreflexivity of \prec_3).
2. Suppose that X and Y are both *Enq* events. Say $E_1 \prec_i E_2 \prec_j E_1$. Relations \prec_2 and \prec_3 do not apply on two *Enq* events, and so we have to check all other possibilities for $i, j \in \{1, 4, 5\}$ when $i \neq j$.
 - (a) Suppose that $\{i, j\} = \{1, 4\}$. Say $E_1 \prec_4 E_2 < E_1$. So there are *Deq* events D_1 and D_2 with $E_1 = \gamma(D_1)$, $E_2 = \gamma(D_2)$ and

$end(D_1) < end(D_2)$. By property 3 of HW applied to $E_2 < \gamma(D_1)$ there is a *Deq* event D_0 such that $E_2 = \gamma(D_0)$ and $end(D_0) < end(D_1)$. That is, $E_2 \prec_4 E_1$, and we are back to case $i = j$.

- (b) Now suppose that $\{i, j\} = \{1, 5\}$. Say $E_1 \prec_5 E_2 < E_1$. So E_1 is in the domain of γ but E_2 is not. But this case is impossible by Lemma 2.4.
 - (c) Finally suppose that $\{i, j\} = \{4, 5\}$. Say $E_1 \prec_4 E_2 \prec_5 E_1$. But this is also impossible since $E_1 \prec_4 E_2$ implies that both of E_1 and E_2 are in the range of γ .
3. Suppose now that one of X and Y is an *Enq* event and the other a *Deq* event: $E \prec D \prec E$. An inspection of relations \prec_i shows that $D \prec E$ implies $D < E$. And $E \prec D$ implies $E < D$ or $E \prec_2 D < E$. Surely it is not the case that $E < D < E$, and so it remains to consider $E \prec_2 D < E$. But $E = \gamma(D)$ would imply $end(E) < end(D)$ by property 2 of the HW list, in contradiction to $D < E$.

q.e.d

Lemma 2.6 *Relation \prec has no cycles.*

Proof. Assume on the contrary that there is a cycle, and take a cycle of minimal length n . In the previous lemma we ruled out the case that $n = 2$, and so we assume that a cycle of minimal length $n > 2$ exists: $X_1 \prec X_2 \cdots \prec X_n \prec X_1$.

- 1. In case all X_i are *Deq* event, we get (with Lemma 2.2) that $end(X_i) < end(X_{i+1})$ and $end(X_n) < end(X_1)$ which is impossible since there are no cycles in the $<$ ordering.
- 2. In case all X_i are *Enq* events there are three possibilities. If no X_i is in the range of γ , then $X_i \prec X_{i+1}$ implies $X_i < X_{i+1}$ and we reach a contradiction since $<$ is an ordering relation. If all X_i 's are in the range of γ then we have a cycle in the \prec_4 relation (Lemma 2.4), which would contradict the transitivity and irreflexivity of \prec_4 . Finally, if some of the X_i 's are in the range of γ and some are not, then we can find two successive *Enq* on the cycle, $E \prec E'$ where E is not in the range and E' is in the range of γ . But then we also have $E' \prec_5 E$ by definition of \prec_5 in contradiction to our earlier result that \prec has no cycles of length two.

3. So we are left with the case that there are both *Enq* and *Deq* events in our minimal cycle. We can find two consecutive members on the cycle a *Deq* event followed by an *Enq* event. Say $X_2 = D$ and $X_3 = E$ form such a pair. Then $D \prec E$ implies $D < E$ since no other relation applies. Consider the possibilities for X_1 :
- (a) If $X_1 = E_1$ is an *Enq* event, then $E_1 \prec D$ allows that either $E_1 < D$ or $E_1 \prec_2 D$. In the first case, $E_1 < D < E$ yields $E_1 < E$ and the cycle can be shortened in contradiction to its minimality. But in the second case, $E_1 = \gamma(D)$ implies $\text{end}(E_1) < \text{end}(D)$ which in light of $D < E$ yields $E_1 < E$ which again shortens the cycle.
 - (b) $X_1 = D_1$ is a *Deq* event. So the situation is that $D_1 \prec D < E$. But $D_1 \prec D$ implies that $D_1 < D$ or $D_1 \prec_3 D$. The first possibility leads again to a shorter cycle, and so we have that $D_1 \prec_3 D$. But this means that $\text{end}(D_1) < \text{end}(D)$, and hence $D_1 < E$ follows which again can be used to shorten the cycle.

q.e.d

So relation \prec has no cycles, and its transitive closure \prec^* is thence an irreflexive ordering relation. Any ordering relation has an extension to a linear ordering and if we extend \prec^* to a linear relation $<^*$ we conclude the proof of theorem 2.1.